

VCP: Single-Stage Cosine Classifier and Prototype Learning for Few-Shot Object Detection

1st Xiang Fei
College of Engineering
Qufu Normal University
Rizhao, China
Fly39@qfnu.edu.cn

2nd Mengru Li
College of Engineering
Qufu Normal University
Rizhao, China
13864085521@qfnu.edu.cn

3rd Jinming Huang*
College of Engineering
Qufu Normal University
Rizhao, China
huangjm@qfnu.edu.cn

4th Jacob Liu
Technology Department
GoerTek Inc.
Weifang, Shandong, China
2374622833@qq.com

5th Xiaofu Zhang
Technology Department
Shandong Free-Optics Technology Co., Ltd.
Weifang, China
fly39@icloud.com

6th Xiaojiao Xu
Technology Department
Shandong Huimi Intelligent Technology Co., Ltd.
Rizhao, China
17604705924@163.com

Abstract—Few-shot object detection (FSOD) methods are almost exclusively built upon the two-stage Faster R-CNN, while YOLO-series single-stage detectors—dominant in industrial deployment—lack systematic FSOD investigation. We present VCP (Visual-Cosine-Prototype), a single-stage FSOD framework combining YOLO11s with a cosine classifier and prototype learning under the TFA protocol. On PASCAL VOC, VCP achieves 72.2% mAP50 (3-split mean) at 10-shot with only 9.5M parameters ($\sim 1/6$ of competing methods). However, the same approach underperforms fine-tuning on MS COCO. Through systematic precision-recall decomposition across 24 comparison groups, we reveal that the cosine classifier acts as a *Recall Booster*—consistently pushing the operating point toward higher recall and lower precision. We formulate the **P-R Tradeoff Hypothesis**: the net benefit depends on the product of baseline recall headroom and the dataset-specific P-R tradeoff ratio, explaining the VOC-COCO divergence. We further explore Florence-2 VLM semantic prototypes as auxiliary priors (+0.5 pp at 10-shot). VCP runs at 54.4 FPS, 4–5 $\times$ faster than two-stage FSOD methods with zero VLM inference overhead.

Index Terms—few-shot object detection, single-stage detector, cosine classifier, prototype learning, recall booster

I. INTRODUCTION

Few-shot object detection (FSOD) aims to detect novel categories from only a few annotated examples. Methods such as FSCE [?] and DeFRCN [?] have steadily advanced FSOD benchmarks, yet a striking commonality persists: virtually all FSOD methods are built upon the two-stage Faster R-CNN [?]. The Region Proposal Network (RPN) provides built-in foreground filtering that helps control false positives under data scarcity. In industrial deployment, however, the YOLO family [?], [?], [?] dominates due to its real-time efficiency. This creates a disconnect where “published methods cannot be

deployed, and deployed systems lack research foundations.” To our knowledge, no prior work has systematically studied modern YOLO detectors (YOLOv8+) for FSOD under the standard TFA protocol.

The core challenge for single-stage FSOD is the absence of RPN foreground filtering: the dense prediction paradigm generates abundant low-quality proposals, subjecting the classifier to extreme imbalance. The cosine classifier [?] mitigates classifier overfitting through normalization, while Florence-2 [?] offers fine-grained semantic understanding. However, combining cosine classification with single-stage YOLO detectors and systematically analyzing its behavior remains unexplored.

We propose VCP (Visual-Cosine-Prototype), which replaces YOLO11s’s linear classification head with a cosine classifier initialized by support-set visual prototypes (Proto mode), and further explores Florence-2 semantic descriptions as auxiliary priors (Fused mode). On PASCAL VOC, VCP achieves 72.2% mAP50 at 10-shot (3-split mean) with ~ 9.5 M parameters, competitive with ~ 60 M-parameter SOTA methods; Split 1 reaches 74.7%, surpassing all existing methods.

However, applying the same approach to MS COCO yields the opposite outcome: cosine-based methods consistently underperform standard fine-tuning. Through systematic P-R decomposition across 24 comparison groups, we identify a consistent pattern: the cosine classifier invariably sacrifices precision for recall—acting as a **Recall Booster**. We formulate the **P-R Tradeoff Hypothesis** to explain this cross-dataset divergence. Our contributions are: (1) the first systematic exploration of single-stage YOLO detectors under the TFA FSOD protocol; (2) identification of the cosine classifier’s Recall Booster property and the P-R Tradeoff Hypothesis; (3) preliminary investigation of VLM semantic prototypes as auxiliary priors and their engineering efficiency.

This work was supported in part by the Key Research and Development Program of Shandong Province, China, under Grant 2025TSGCXTHG031, and in part by the National Key Research and Development Program of China under Grant 2021YFE0193900. (Corresponding author: Jinming Huang.) Jinming Huang is with the College of Engineering, Qufu Normal University, Shandong 276826, China (email: huangjm@qfnu.edu.cn).

II. RELATED WORK

FSOD research has transitioned from meta-learning approaches [?], [?] to the fine-tuning paradigm established by TFA [?], which trains a detector on base classes then fine-tunes the classification head on novel classes. FSCE [?] introduced contrastive proposal encoding, DeFRCN [?] proposed gradient decoupling to mitigate backbone overfitting, and KD-DeFRCN [?] incorporated knowledge distillation. Other improvements include MPSR [?] for multi-scale refinement, FSDetView [?] for viewpoint augmentation, and CME [?] for margin optimization. All these methods share a common backbone: Faster R-CNN, exploiting its RPN for proposal quality control.

Single-stage detector exploration in FSOD remains extremely limited. Meta-YOLO [?] modulated YOLOv2 features via meta-learning, but used a high-variance split protocol since superseded by TFA’s low-variance splits [?]. No prior work has systematically studied modern YOLO detectors (v8+) under the TFA protocol. In parallel, VLMs such as CLIP [?] and Florence-2 [?] have enabled semantic prototype construction—directly obtaining class priors from textual descriptions rather than scarce visual samples—while OWL-ViT [?], GLIP [?], and Grounding DINO [?] have advanced open-vocabulary detection. However, leveraging VLM prototypes under FSOD’s strict data constraints (1–10 samples per class) without introducing excessive semantic noise remains underexplored.

III. METHOD

A. Framework and Prototype Construction

VCP adopts YOLO11s as the base detector with a critical modification: the linear classification head is replaced by cosine similarity classification (Fig. ??). For each proposal feature $f \in \mathbb{R}^d$, class scores are computed as $s_c = \cos(f, p_c) / \tau$, where p_c is the prototype for class c and τ a learnable temperature. The normalization $\|f\| = \|p_c\| = 1$ makes classification depend solely on directional consistency, mitigating overfitting.

Prototypes are derived from two sources. **Visual Prototype (Proto):** For each novel class c , RoIAlign extracts features within GT boxes from FPN P3–P5 at the classification branch, averaged and L2-normalized: $p_c^{\text{vis}} = \frac{1}{K} \sum_{i=1}^K \text{RoIAlign}(F_i, b_i)$, written directly into classifier weights.

Semantic Prototype (Fused): Target crops are fed to Florence-2 with <DETAILED_CAPTION>, text descriptions encoded via BART to p_c^{text} . The fused prototype is:

$$p_c = \alpha \cdot p_c^{\text{vis}} + (1 - \alpha) \cdot W_f \cdot p_c^{\text{text}} \quad (1)$$

where $\alpha = 0.7$ and W_f is a learnable projection matrix. Critically, semantic prototype extraction occurs only during training; at inference, prototypes are fixed in the classifier weights, incurring zero VLM overhead.

B. Cosine Classifier as a Recall Booster

The gradient of cosine similarity w.r.t. feature f reveals the mechanism behind its behavioral pattern:

$$\frac{\partial \cos(f, p_c)}{\partial f} = \frac{p_c - \cos(f, p_c) \cdot \hat{f}}{\|f\|}, \quad \hat{f} = \frac{f}{\|f\|} \quad (2)$$

When f has low similarity to p_c , the gradient magnitude is large, driving the decision boundary outward. When similarity is already high, gradients diminish. Unlike linear classifiers whose gradient $\partial(W_c f) / \partial f = W_c$ is uniform, the cosine classifier *disproportionately updates low-confidence samples*, naturally favoring a high-recall detection strategy: it detects more true positives (high recall) but admits more false positives (low precision). We term this property the **Recall Booster**. Whether this yields a net mAP gain depends on the product of baseline recall headroom and the dataset’s P-R tradeoff ratio, as validated in Section IV.

IV. EXPERIMENTS

A. Setup

We evaluate on PASCAL VOC 2007+2012 and MS COCO 2014 under the TFA split protocol [?]: VOC’s 20 classes partitioned into 15 base / 5 novel (3 splits); COCO’s 20 VOC-overlapping classes as novel, 60 as base. Base detector is YOLO11s (9.5M params), pre-trained on COCO base classes. Stage 1: 100 epochs on base classes, batch 64, lr 0.01. Stage 2: fine-tune on novel classes up to 200 epochs with early stopping (patience 30), batch 4, lr 0.0005. SGD optimizer (momentum 0.937, weight decay 0.0005, 3-epoch warmup), backbone frozen first 10 epochs. Data augmentation: mosaic, RandAugment, HSV jitter, random flip/scale/translate, and random erasing. Input size 640×640 . Temperature τ initialized at 5.0 (learnable). All experiments on a single NVIDIA H20 GPU.

Four variants are compared: (1) *Finetune*—standard YOLO11s with linear classifier; (2) *Cosine*—cosine classifier, random initialization; (3) *Cosine+Proto*—cosine classifier initialized with visual prototypes; (4) *Cosine+Fused*—cosine classifier with visual + Florence-2 fused prototypes.

B. Main Results

VOC Results. Table ?? presents VOC mAP50 (Split 1 / 3-split mean $\pm$ std). We observe three trends: (1) The cosine classifier drives a qualitative leap—5-shot (Split 1): 32.9% $\rightarrow$ 66.9% (+34.0 pp); 10-shot (Split 1): 51.3% $\rightarrow$ 70.7% (+19.4 pp). (2) Visual prototype initialization contributes +3.7 pp at 10-shot (Split 1). (3) Florence-2 fusion adds +0.3 pp at 10-shot (Split 1), directionally consistent across all three splits.

COCO Results. Table ?? shows a contrasting pattern: Fine-tune outperforms all cosine variants. At 30-shot, Cosine+Fused reaches 52.3%, improving over pure Cosine (49.4%) but still below Finetune (60.3%). Results are on a single split and require further verification.

VCP (VLM-Cosine-Prototype) Architecture for Few-Shot Object Detection




Fig. 1. VCP framework: YOLO11s backbone+neck, decoupled head, cosine classifier. Visual prototypes via RoIAlign; semantic prototypes via Florence-2 + BART. Fused at $\alpha = 0.7$.

TABLE I
VOC MAP50 (SPLIT 1 / 3-SPLIT MEAN $\pm$ STD)

Method	1-shot	3-shot	5-shot	10-shot
Finetune	10.9 / -	20.5 / -	32.9 / -	51.3 / -
Cosine	12.8 / -	51.3 / -	66.9 / 57.8 $\pm$ 2.9	70.7 / 69.0 $\pm$ 3.1
Cosine+Proto	21.5 / -	52.3 / -	68.8 / 57.8 $\pm$ 2.8	74.4 / 71.7 $\pm$ 2.7
Cosine+Fused	21.3 / -	53.9 / -	69.0 / 58.1 $\pm$ 2.5	74.7 / 72.2 $\pm$ 2.6

TABLE II
COCO MAP50

Method	10-shot	30-shot
Finetune	59.0	60.3
Cosine	47.1	49.4
Cosine+Proto	47.2	51.6
Cosine+Fused	47.0	52.3

TABLE III
P-R DECOMPOSITION SUMMARY

Setting	ΔP	ΔR	ΔmAP	$\Delta R/ \Delta P $
VOC 5-shot Cos+Fused	-0.22	+0.45	+35.9	2.00
VOC 10-shot Cos+Fused	-0.35	+0.33	+23.4	0.96
COCO 10-shot Cos+Fused	-0.81	+0.39	-12.0	0.48
COCO 30-shot Cos+Fused	-0.56	+0.14	-8.0	0.25

C. P-R Analysis and the P-R Tradeoff Hypothesis

To explain the VOC-COCO divergence, we decompose Precision and Recall across all 24 method $\times$ shot $\times$ dataset combinations. Table ?? summarizes the key patterns.

Three findings are observed (vs. Finetune baseline):

Finding 1—Invariant Recall Booster: All 24 cosine variants *without exception* shift toward higher recall and lower

precision—a systematic property (see gradient analysis, Section III-B).

Finding 2—Tradeoff Ratio Determines Net Benefit: On VOC, $\Delta R/|\Delta P| \approx 2.0$ at 5-shot; recall gain dominates. On COCO, the ratio collapses to 0.48 (10-shot) and 0.25 (30-shot), with ΔP far exceeding ΔR in magnitude. COCO’s 20 novel classes (vs. 5 on VOC) cause massive cross-class false

positives under the relaxed decision boundary.

Finding 3—Fused Mode Modulates P-R Balance: VOC 10-shot Fused ($P=0.336$, $R=0.848$) recovers precision (+0.136 vs. pure Cosine at $P=0.200$, $R=0.872$) with minimal recall loss (-0.024), suggesting semantic priors moderate prototype over-relaxation.

We formalize these observations as the **P-R Tradeoff Hypothesis**: the cosine classifier inherently pushes toward high-recall, low-precision operation; net benefit is the product of (1) baseline recall headroom and (2) the dataset P-R tradeoff ratio, determined by class count and scene complexity.

D. Ablation and SOTA Comparison

Ablation. Table ?? merges component and prototype source ablations. The cosine head is the dominant driver ($\sim 90\%$ of gains). Visual prototypes alone (74.4%) outperform Florence-2 text-only (72.1%), and fusion achieves the optimum (74.7%), confirming complementary value.

TABLE IV
ABLATION: COMPONENT & PROTOTYPE SOURCE (VOC)

Configuration	3-shot	5-shot	10-shot
Finetune (baseline)	20.5	32.9	51.3
+ Cosine head	51.3	66.9	70.7
+ Visual Proto	52.3	68.8	74.4
+ VLM Fused	53.9	69.0	74.7
Florence-2 Only (10-shot)	—	—	72.1

SOTA Comparison. Table ?? compares VCP against existing FSOD methods (VOC Split 1). All baselines use Faster R-CNN R-101 ($\sim 60M$), VCP uses YOLO11s ($\sim 9.5M$). At 5-shot and 10-shot, VCP matches or surpasses all prior methods while using $\sim 1/6$ the parameters. The 1-shot gap indicates single-stage detectors need additional regularization under extreme scarcity.

TABLE V
COMPARISON WITH SOTA (VOC SPLIT 1, MAP50)

Method	1-shot	3-shot	5-shot	10-shot
TFA w/cos	39.8	44.7	55.7	56.0
FSCE	44.2	51.4	61.9	63.4
DeFRCN	53.6	61.5	64.1	60.8
KD-DeFRCN	58.2	65.1	68.2	67.4
VCP Cos+Fused	21.3	53.9	69.0	74.7
VCP Cos+Proto	21.5	52.3	68.8	74.4

E. Inference Speed

Table ?? compares inference speed (NVIDIA H20, 640×640 , batch 1). VCP Fused achieves 54.4 FPS—4–5 $\times$ faster than two-stage FSOD methods—with only 9.5M parameters. The cosine classifier adds ~ 4 ms latency (from L2 normalization), still well above the 30 FPS real-time threshold. Critically, Florence-2 is *not* loaded at inference; semantic prototypes are pre-computed and stored in the classifier weights.

TABLE VI
INFERENCE SPEED (H20, 640×640 , BS=1)

Model	Latency (ms)	FPS	Params
FRCN R-101 (ref.)	~ 85	~ 12	$\sim 60M$
YOLO11s baseline	14.4	69.5	9.5M
Finetune	14.7	68.2	9.5M
Cosine	18.9	52.8	9.5M
Fused (VCP)	18.4	54.4	9.5M

V. CONCLUSION

We presented VCP, a single-stage FSOD framework combining YOLO11s with cosine classification and prototype learning. Key findings: (1) Single-stage FSOD is viable—VCP achieves competitive performance at $1/6$ the parameters of two-stage methods, at 54.4 FPS. (2) Through 24-group P-R decomposition, we identified the Recall Booster property and P-R Tradeoff Hypothesis, explaining VOC-COCO divergence via baseline recall headroom and dataset tradeoff ratio. (3) Florence-2 semantic prototypes provide consistent but limited gains with zero inference overhead.

Limitations include weak 1-shot performance, single-split COCO evaluation, and the post-hoc nature of the tradeoff hypothesis. Future work will explore adaptive VLM fusion, cross-domain generalization, and independent hypothesis validation.

REFERENCES

- [1] B. Sun et al., “FSCE: Few-shot object detection via contrastive proposal encoding,” in *CVPR*, 2021, pp. 7352–7362.
- [2] L. Qiao et al., “DeFRCN: Decoupled faster R-CNN for few-shot object detection,” in *ICCV*, 2021, pp. 8681–8690.
- [3] S. Ren et al., “Faster R-CNN: Towards real-time object detection with region proposal networks,” in *NeurIPS*, 2015.
- [4] J. Redmon et al., “You only look once: Unified, real-time object detection,” in *CVPR*, 2016, pp. 779–788.
- [5] G. Jocher et al., “Ultralytics YOLO (version 8.0),” 2023. [Online]. Available: <https://github.com/ultralytics/ultralytics>
- [6] X. Wang et al., “Frustratingly simple few-shot object detection,” in *ICML*, 2020, pp. 9919–9928.
- [7] A. Radford et al., “Learning transferable visual models from natural language supervision,” in *ICML*, 2021, pp. 8748–8763.
- [8] B. Xiao et al., “Florence-2: Advancing a unified representation for a variety of vision tasks,” in *CVPR*, 2024, pp. 4818–4829.
- [9] X. Yan et al., “Meta R-CNN: Towards general solver for instance-level low-shot learning,” in *ICCV*, 2019, pp. 9577–9586.
- [10] Y.-X. Wang et al., “Meta-learning to detect rare objects,” in *ICCV*, 2019, pp. 9925–9934.
- [11] W. Pei et al., “Few-shot object detection by knowledge distillation using bag-of-visual-words representations,” in *ECCV*, 2022, pp. 283–299.
- [12] J. Wu et al., “Multi-scale positive sample refinement for few-shot object detection,” in *ECCV*, 2020, pp. 456–472.
- [13] Y. Xiao and R. Marlet, “Few-shot object detection and viewpoint estimation for objects in the wild,” in *ECCV*, 2020, pp. 192–210.
- [14] B. Li et al., “Beyond max-margin: Class margin equilibrium for few-shot object detection,” in *CVPR*, 2021, pp. 7363–7372.
- [15] J. Redmon and A. Farhadi, “YOLOv3: An incremental improvement,” arXiv:1804.02767, 2018.
- [16] B. Kang et al., “Few-shot object detection via feature reweighting,” in *ICCV*, 2019, pp. 8420–8429.
- [17] M. Minderer et al., “Simple open-vocabulary object detection with vision transformers,” in *ECCV*, 2022.
- [18] L. H. Li et al., “Grounded language-image pre-training,” in *CVPR*, 2022, pp. 10965–10975.

- [19] S. Liu et al., “Grounding DINO: Marrying DINO with grounded pre-training for open-set object detection,” arXiv:2303.05499, 2023.